



# kubernetes

## Content List

Kubernetes (K8s) .....	2
What is kubernetes? .....	2
Features of Kubernetes .....	3
Kubernetes Architecture.....	4
Here is an overview of the main components of Kubernetes architecture .....	4
Install Kubernetes .....	6
Multi-Node Kubernetes clusters .....	7
Overview of Setting up a Multi-node Kubernetes Cluster... <b>Error! Bookmark not defined.</b>	
Kubernetes Setup Using Kubeadm In AWS EC2 Ubuntu Servers Container-D As Runtime	8
Add Worker Machines to Kubernetes Master.....	19
Deploy Sample Application .....	23
How to deploy/run/execute tasks/workloads in kubernetes .....	24
Declarative Approach.....	25
More about Pods .....	27
Pod Management Using Controllers.....	28
Pod Lifecycle and Communication.....	32
How do pods communicate with each other? .....	33
Some commonly used kubectl commands.....	36
Differences between ReplicaSets and Replication Controller .....	39
Example: Deploy a nodeApp in a ReplicaSet.....	40

DaemonSet:.....	42
Example: Deploy a log management (logmgt) in a DaemonSet.....	42
Deployment controller .....	44
Example: Deployment of a hello app in a deployment using Recreate strategy.....	44
Deployment techniques:.....	46
Pod Auto-Scaling .....	49
How to check resource usage in Kubernetes .....	50
Installation of Metric Server .....	51
Deployment with HPA .....	52
Example: Deployment of Spring app and MongoDB as pod without volumes.....	53
Kubernetes Volumes: .....	54
Configuration of NFS Server .....	55
Configuring the NFS Client Machine .....	56
Deployment of a MongoDB using the NFS.....	56
Setting up Kubernetes cluster on AWS using KOPS .....	57
Persistent Volumes .....	60
Kubernetes ConfigMaps and Secrets .....	66

## **Kubernetes (K8s)**

### **What is kubernetes?**

- Kubernetes is an orchestration engine and open-source platform for managing containerized applications.
- Responsibilities include container deployment, scaling and descaling of containers and container load balancing.
- Actually, kubernetes is not a replacement for Docker, but kubernetes can be considered as a replacement for Docker swarm, kubernetes is significantly more complex than docker swarm and requires more work to deploy.
- Born in Google, written in Go/Golang. Donated to CNCF (Cloud Native Computing Foundation) in 2014.

- Kubernetes v1.0 was released on July 21, 2015.
- Current stable release v1.30.0.

### Features of Kubernetes

1. **Automated Deployment and Scaling:** Kubernetes allows you to define desired states for your applications and then automates the process of deploying, updating, and scaling them. It can automatically roll out and roll back changes based on defined criteria.
2. **Service Discovery and Load Balancing:** Kubernetes provides built-in service discovery and load balancing, allowing you to expose your applications as network services and balance the load across different instances.
3. **Self-Healing:** Kubernetes continuously monitors the health of your applications and can automatically restart containers that fail, replace containers, or reschedule them on other nodes if necessary.
4. **Horizontal Scaling:** Kubernetes supports horizontal scaling of applications, allowing you to add or remove instances of your applications dynamically based on demand.
5. **Storage Orchestration:** Kubernetes manages persistent storage using various storage systems and providers, enabling you to dynamically provision, manage, and attach storage to your applications.
6. **Configuration Management:** Kubernetes supports dynamic configuration and secrets management through ConfigMaps and Secrets, allowing you to separate application configuration from the code and manage it more efficiently.
7. **Declarative API:** Kubernetes uses a declarative API, allowing you to specify the desired state of your applications and letting Kubernetes handle the complexities of achieving that state.
8. **Namespace Isolation:** Kubernetes supports multi-tenancy through namespaces, which allow you to create isolated environments for different teams or projects within a single cluster.
9. **Resource Management:** Kubernetes provides mechanisms for managing and optimizing the allocation of resources such as CPU and memory across different applications, ensuring efficient utilization.

10. **Built-in Monitoring and Logging:** Kubernetes offers integration with monitoring and logging tools to track the health and performance of your applications.
11. **Role-Based Access Control (RBAC):** Kubernetes provides fine-grained access control to your cluster resources through RBAC, allowing you to manage user permissions and access levels.
12. **Helm:** While not a core feature of Kubernetes, Helm is a package manager for Kubernetes that simplifies the management of applications and services within the cluster.

### **Kubernetes Architecture**

Kubernetes architecture is designed to manage and orchestrate containerized applications in a cluster of nodes.

It follows a client-server model and consists of several components that work together to achieve this.

**Here is an overview of the main components of Kubernetes architecture:**

1. **Cluster:**
  - A Kubernetes cluster consists of a set of worker nodes and at least one control plane node (master node).
2. **Control Plane (Master Node):**
  - The control plane is responsible for managing the cluster and orchestrating the deployment of applications. It includes the following components:
    - **kube-apiserver:** The API server is the entry point to the Kubernetes cluster and exposes the Kubernetes API. It handles all API requests and is the primary interface for interacting with the cluster.
    - **kube-scheduler:** The scheduler is responsible for determining which node a new pod should be scheduled on based on resource availability and other constraints.
    - **kube-controller-manager:** The controller manager runs various controllers that watch for changes in the cluster state and take corrective actions as needed, such as replicating pods or handling node failures.

- **cloud-controller-manager:** This component manages interactions with the underlying cloud provider's infrastructure (if applicable) for tasks such as load balancing and persistent storage.
- **etcd:** etcd is a distributed key-value store that stores the cluster's state, including configuration and metadata. It serves as the source of truth for the cluster.

### 3. Worker Nodes:

- Worker nodes run the applications (pods) and include the following components:
  - **kubelet:** The kubelet is an agent that runs on each node. It ensures that the containers described in pod specs are running and manages their lifecycle.
  - **kube-proxy:** The kube-proxy manages network rules on the node, allowing network communication to and from pods.
  - **Container Runtime:** The container runtime is responsible for running and managing containers on the node. Examples include Docker and containerd.

### 4. Pods:

- Pods are the smallest deployable units in Kubernetes and represent a group of one or more containers that share network and storage resources.

### 5. Services:

- Services are a way to expose pods and provide stable network endpoints for them. They can also handle load balancing.

### 6. Networking:

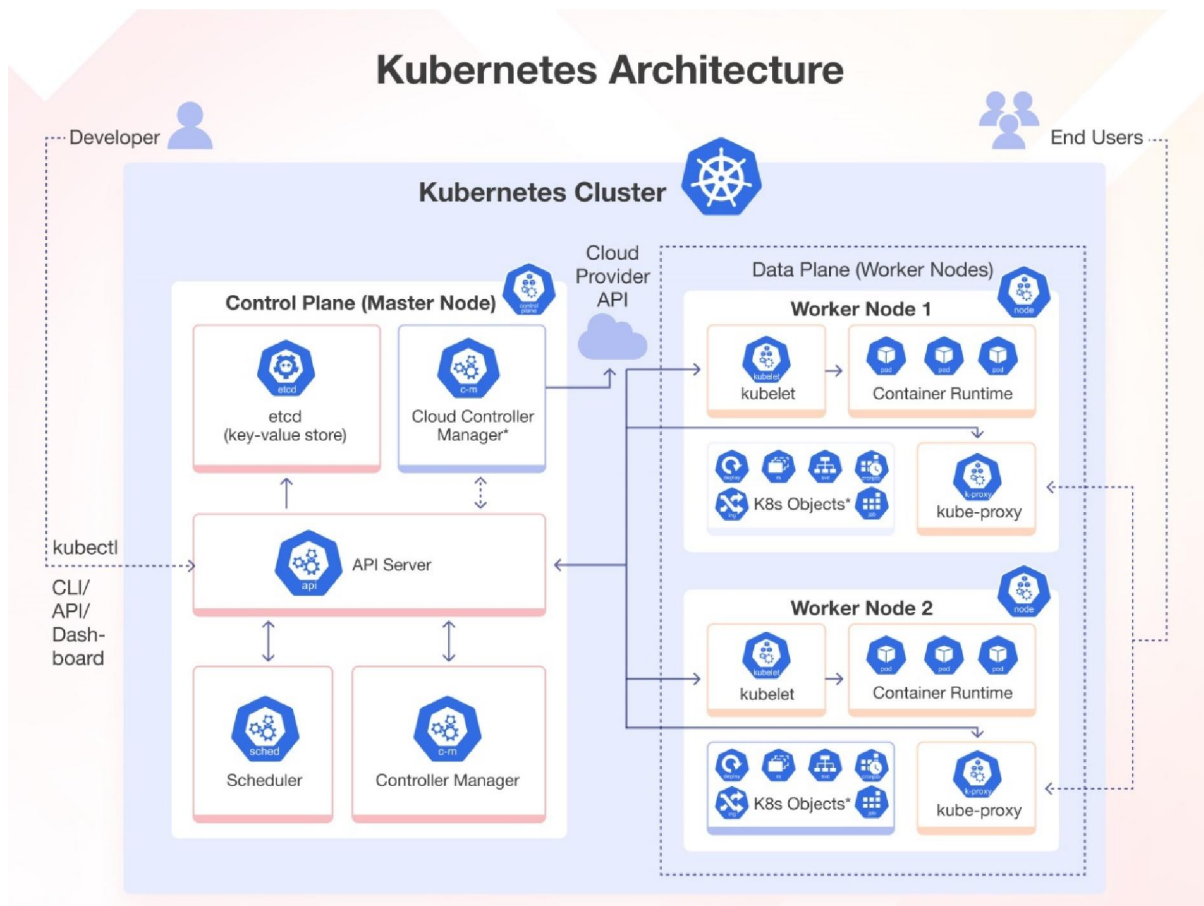
- Kubernetes provides a flat network model that allows all pods to communicate with each other without Network Address Translation (NAT).

### 7. Ingress:

- Ingress is a resource that manages external access to services, providing load balancing, SSL termination, and path-based routing.

## 8. Namespaces:

- Namespaces provide a way to divide cluster resources among different teams or projects, allowing for resource isolation and management.



## Install Kubernetes

Installing Kubernetes on a single system is mainly for development, testing, and learning purposes. There are several tools available that allow you to set up a single-node Kubernetes cluster quickly and easily. Here is an example of the most popular options:

### Minikube

- **Description:** Minikube is a tool that enables you to run a single-node Kubernetes cluster locally on your machine. It's suitable for development and testing.
- **Steps:**

1. **Install Minikube:** Download and install Minikube on your system from the official website.
2. **Install kubectl:** Ensure **kubectl** is installed to interact with the Kubernetes cluster.
3. **Start Minikube:** Run **minikube start** to start the cluster.
4. **Configure kubectl:** Minikube automatically configures **kubectl** to connect to the local cluster.
5. **Verify the installation:** Use commands like **kubectl get nodes** to verify the cluster is running.

### Multi-Node Kubernetes clusters

1. Self-managed kubernetes cluster

kubeadm → We can setup multi-node k8's cluster using kubeadm.

kubespary → We can setup multi-node k8s cluster using kubespary (Ansible playbooks used internally by kubespary).

It needs manually management of the control plane and worker nodes.

2. Managed k8s cluster (Cloud services)

The control plane and its components are managed by the cloud provider. Examples:

- ✓ EKS → Elastic Kubernetes Service (AWS)
- ✓ AKS → Azure Kubernetes Service (Azure)
- ✓ GKE → Google Kubernetes Engine (GCP)
- ✓ IKE → IBM Kubernetes Engine (IBM Cloud)

3. **KOPs (Kubernetes Operations)**

It is a software used to create a production ready kubernetes in AWS.

It is highly available k8s services in cloud like AWS.

It will leverage cloud services like:

- Autoscaling groups
- Load balancer
- Launch template/configuration
- Nodes (WorkerNodes and MasterNodes)

4. **Rancher**

With Rancher, we can deploy both managed and self-managed k8s

It serves as a glass to access and manage multiple k8s from the UI- Rancher Dashboard.

Rancher needs authentication and authorization access the different cloud platforms.

## Kubernetes Setup Using Kubeadm in AWS EC2 Ubuntu Servers Container-D As Runtime

### **Prerequisite:**

### **3 - Ubuntu Servers**

- ✓ 1 - Manager (4GB RAM , 2 Core) t2.medium
- ✓ 2 - Workers (1 GB, 1 Core) t2.micro
- ✓ **Note:** Open Required Ports in AWS Security Groups. For now, we will open All traffic.

### **=====IMPERATIVE APPROACH TO SETUP KUBERNETES IN AWS=====**

1. Assign hostname & login as 'root' user because the following set of commands need to be executed with 'sudo' permissions

```
sudo hostnamectl set-hostname master
sudo -i
```

```
# run the following below as a script
# This will Install Required packages and apt keys.
#!/bin/bash
sudo apt update -y
sudo apt install -y apt-transport-https
sudo curl -s
https://packages.cloud.google.com/apt/doc/apt-key.gpg |
apt-key add -
cat <<EOF >/etc/apt/sources.list.d/kubernetes.list
deb https://apt.kubernetes.io/ kubernetes-xenial main
EOF
apt update -y
#Turn Off Swap Space
swapoff -a
sed -i '/ swap / s/^\(.*\)$/#\1/g' /etc/fstab
```



```

#Install kubeadm, Kubelet And Kubectl containerd
sudo apt-get install -y kubelet containerd kubeadm kubectl
kubernetes-cni
# apt-mark hold will prvent the package from being
authomatically upgraded or removed
sudo apt-mark hold kubelet containerd kubeadm kubectl
kubernetes-cni
#
cat <<EOF | sudo tee /etc/modules-load.d/containerd.conf
overlay
br_netfilter
EOF

# Setup required sysctl params, these persist across
reboots.
cat <<EOF | sudo tee /etc/sysctl.d/99-kubernetes-cri.conf
net.bridge.bridge-nf-call-iptables = 1
net.ipv4.ip_forward = 1
net.bridge.bridge-nf-call-ip6tables = 1
EOF

# Apply sysctl params without reboot
sudo sysctl --system
# Configure containerd:
sudo mkdir -p /etc/containerd
containerd config default | sudo tee
/etc/containerd/config.toml
# Restart containerd:
sudo systemctl restart containerd
# If you get error releated to kubernetes-cni if alreay
exists install with out kubernetes-cni
apt-get install -y kubelet kubeadm kubectl
# Enable and start kubelet service
sudo systemctl daemon-reload

```

```
sudo systemctl start kubelet  
sudo systemctl enable kubelet.service
```

2. Exit as root user & execute the below commands as normal ubuntu user

```
sudo su - ubuntu
```

3. Initialised the control plane

```
# Initialize Kubernetes master by executing below  
command.  
sudo kubeadm init  
  
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

4. To verify, if kubectl is working or not, run the following command

```
kubectl get pods -A
```

5. Deploy the network plugin - weave network

```
#deploy the network plugin - weave network  
kubectl apply -f  
https://github.com/weaveworks/weave/releases/download/v2  
.8.1/weave-daemonset-k8s.yaml  
kubectl get pods -A  
kubectl get node
```

6. Copy kubeadm join token from the master and execute in Worker Nodes to join to cluster

```
kubeadm join 172.31.10.12:6443 --token
cdm6fo.dhbrxyleqe5suy6e \
--discovery-token-ca-cert-hash
sha256:1fc51686afd16c46102c018acb71ef9537c1226e331840e7d
401630b96298e7d
```

## =====SCRIPT TO SETUP KUBERNETES, COMMON FOR MASTER & SLAVES

=====START=====

```
#!/bin/bash
#1) Switch to root user [ sudo -i]

sudo hostnamectl set-hostname node2
sudo -i

#2) Disable swap & add kernel settings

swapoff -a
sed -i 's/ swap / s/^\(.*\)$/#\1/g' /etc/fstab

#3) Add kernel settings & Enable IP tables(CNI Prerequisites)

cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
overlay
br_netfilter
EOF

modprobe overlay
modprobe br_netfilter

cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-ip6tables = 1
net.ipv4.ip_forward = 1
EOF

sysctl --system

#4) Install containerd run time

#To install containerd, first install its dependencies.
```

```

apt-get update -y
apt-get install ca-certificates curl gnupg lsb-release -y

#Note: We are not installing Docker Here. Since containerd.io package is
part of docker apt repositories hence we added docker repository & it's
key to download and install containerd.
# Add Docker's official GPG key:
sudo mkdir -p /etc/apt/keyrings
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --
dearmor -o /etc/apt/keyrings/docker.gpg

#Use following command to set up the repository:

echo \
  "deb [arch=$(dpkg --print-architecture) signed-
by=/etc/apt/keyrings/docker.gpg] https://download.docker.com/linux/ubuntu
  \
  $(lsb_release -cs) stable" | sudo tee
/etc/apt/sources.list.d/docker.list > /dev/null

# Install containerd

apt-get update -y
apt-get install containerd.io -y

# Generate default configuration file for containerd

#Note: Containerd uses a configuration file located in
/etc/containerd/config.toml for specifying daemon level options.
#The default configuration can be generated via below command.

containerd config default > /etc/containerd/config.toml

# Run following command to update configure cgroup as systemd for
containerd.

sed -i 's/SystemdCgroup \= false/SystemdCgroup \= true/g'
/etc/containerd/config.toml

# Restart and enable containerd service

systemctl restart containerd
systemctl enable containerd

#5) Installing kubeadm, kubelet and kubectl

# Update the apt package index and install packages needed to use the
Kubernetes apt repository:

apt-get update
apt-get install -y apt-transport-https ca-certificates curl

# Download the Google Cloud public signing key:

curl -fsSL https://packages.cloud.google.com/apt/doc/apt-key.gpg

```

```

# Add the Kubernetes apt repository:

echo "deb [signed-by=/etc/apt/keyrings/kubernetes-archive-keyring.gpg]
https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee
/etc/apt/sources.list.d/kubernetes.list

# Update apt package index, install kubelet, kubeadm and kubectl, and pin
their version:

# Add Kubernetes APT keyring and repository
sudo curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.28/deb/Release.key |
sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
echo "deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.28/deb/ /" | sudo tee
/etc/apt/sources.list.d/kubernetes.list > /dev/null

# Update package index
sudo apt update

# Install Kubernetes packages
sudo apt install -y kubelet kubeadm kubectl

# apt-mark hold will prevent the package from being automatically
upgraded or removed.

apt-mark hold kubelet kubeadm kubectl

# Enable and start kubelet service

systemctl daemon-reload
systemctl start kubelet
systemctl enable kubelet.service

#Copy kubeadm join token from the master and execute in Worker Nodes to
join to cluster
# replace this token with yours

kubeadm join 172.31.63.97:6443 --token hapn6z.ufwr4b82282z3aix \
--discovery-token-ca-cert-hash
sha256:d45cee841b1694ff440baabf11067b2b7c237bc60b1df5eab4983f4dbae6011c

```

===== **END** =====

=====In Master Node Only=====

#kubeadm command in the master node will generate the config for the setup of the entire cluster.

```
kubeadm init
```

## =====Start of Master Node Script=====

```
#!/bin/bash
#1) Switch to root user [ sudo -i]

sudo hostnamectl set-hostname node1

#2) Disable swap & add kernel settings

swapoff -a
sed -i '/ swap / s/^\(.*\)$/#\1/g' /etc/fstab

#3) Add kernel settings & Enable IP tables(CNI Prerequisites)

cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
overlay
br_netfilter
EOF

modprobe overlay
modprobe br_netfilter

cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-ip6tables = 1
net.ipv4.ip_forward = 1
EOF

sysctl --system

#4) Install containerd run time
```

```

#To install containerd, first install its dependencies.

apt-get update -y
apt-get install ca-certificates curl gnupg lsb-release -y

#Note: We are not installing Docker Here. Since containerd.io
package is part of docker apt repositories hence we added
docker repository & it's key to download and install
containerd.

# Add Docker's official GPG key:
sudo mkdir -p /etc/apt/keyrings
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo
gpg --dearmor -o /etc/apt/keyrings/docker.gpg

#Use following command to set up the repository:

echo \
    "deb [arch=$(dpkg --print-architecture) signed-
by=/etc/apt/keyrings/docker.gpg]
https://download.docker.com/linux/ubuntu \
    $(lsb_release -cs) stable" | sudo tee
/etc/apt/sources.list.d/docker.list > /dev/null

# Install containerd

apt-get update -y
apt-get install containerd.io -y

# Generate default configuration file for containerd

#Note: Containerd uses a configuration file located in
/etc/containerd/config.toml for #specifying daemon level
options.

#The default configuration can be generated via below command.

```



```
containerd config default > /etc/containerd/config.toml

# Run following command to update configure cgroup as systemd
for containerd.

sed -i 's/SystemdCgroup \= false/SystemdCgroup \= true/g'
/etc/containerd/config.toml

# Restart and enable containerd service

systemctl restart containerd
systemctl enable containerd

#5) Installing kubeadm, kubelet and kubectl

# Update the apt package index and install packages needed to
use the Kubernetes apt repository:

apt-get update
apt-get install -y apt-transport-https ca-certificates curl

# Download the Google Cloud public signing key:

curl -fsSLo /etc/apt/keyrings/kubernetes-archive-keyring.gpg
https://packages.cloud.google.com/apt/doc/apt-key.gpg

# Add the Kubernetes apt repository:

echo "deb [signed-by=/etc/apt/keyrings/kubernetes-archive-
keyring.gpg] https://apt.kubernetes.io/ kubernetes-xenial
main" | sudo tee /etc/apt/sources.list.d/kubernetes.list
```

```
# Update apt package index, install kubelet, kubeadm and
kubectl, and pin their version:

apt-get update
apt-get install -y kubelet kubeadm kubectl

# apt-mark hold will prevent the package from being
automatically upgraded or removed.

apt-mark hold kubelet kubeadm kubectl

# Enable and start kubelet service

systemctl daemon-reload
systemctl start kubelet
systemctl enable kubelet.service

# initialise the control plane
kubeadm init

# Configure kubectl as a normal ubuntu user in a different
node to manage #to manage the cluster from it and not from
master node. This is for #security reasons
sudo su - ubuntu

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

# To verify, if kubectl is working or not, run the following
command.

kubectl get pods -o wide -n kube-system

#You will notice from the previous command, that all the pods
are running except one: 'core-dns'. For resolving this we
```



```
cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
```

```
overlay
```

```
br_netfilter
```

```
EOF
```

```
modprobe overlay
```

```
modprobe br_netfilter
```

```
cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
```

```
net.bridge.bridge-nf-call-iptables = 1
```

```
net.bridge.bridge-nf-call-ip6tables = 1
```

```
net.ipv4.ip_forward = 1
```

```
EOF
```

```
sysctl --system
```

#4) Install containerd run time

#To install containerd, first install its dependencies.

```
apt-get update -y
```

```
apt-get install ca-certificates curl gnupg lsb-release -y
```

#Note: We are not installing Docker Here. Since containerd.io package is part of docker apt repositories hence we added docker repository & it's key to download and install containerd.

# Add Docker's official GPG key:

```
sudo mkdir -p /etc/apt/keyrings
```

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo  
gpg --dearmor -o /etc/apt/keyrings/docker.gpg
```

#Use following command to set up the repository:

```

echo \
    "deb [arch=$(dpkg --print-architecture) signed-
by=/etc/apt/keyrings/docker.gpg]
https://download.docker.com/linux/ubuntu \
    $(lsb_release -cs) stable" | sudo tee
/etc/apt/sources.list.d/docker.list > /dev/null

# Install containerd

apt-get update -y
apt-get install containerd.io -y

# Generate default configuration file for containerd

#Note: Containerd uses a configuration file located in
/etc/containerd/config.toml for specifying daemon level
options.
#The default configuration can be generated via below command.

containerd config default > /etc/containerd/config.toml

# Run following command to update configure cgroup as systemd
for containerd.

sed -i 's/SystemdCgroup \= false/SystemdCgroup \= true/g'
/etc/containerd/config.toml

# Restart and enable containerd service

systemctl restart containerd
systemctl enable containerd

#5) Installing kubeadm, kubelet and kubect1

```

```
# Update the apt package index and install packages needed to
use the Kubernetes apt repository:

apt-get update
apt-get install -y apt-transport-https ca-certificates curl

# Download the Google Cloud public signing key:

curl -fsSLo /etc/apt/keyrings/kubernetes-archive-keyring.gpg
https://packages.cloud.google.com/apt/doc/apt-key.gpg

# Add the Kubernetes apt repository:

echo "deb [signed-by=/etc/apt/keyrings/kubernetes-archive-
keyring.gpg] https://apt.kubernetes.io/ kubernetes-xenial
main" | sudo tee /etc/apt/sources.list.d/kubernetes.list

# Update apt package index, install kubelet, kubeadm and
kubectl, and pin their version:

apt-get update
apt-get install -y kubelet kubeadm kubectl

# apt-mark hold will prevent the package from being
automatically upgraded or removed.

apt-mark hold kubelet kubeadm kubectl

# Enable and start kubelet service

systemctl daemon-reload
systemctl start kubelet
systemctl enable kubelet.service
```

```

#Copy kubeadm join token from the master and execute in Worker
Nodes to join to cluster
# replace this token with yours

kubeadm join 10.0.0.11:6443 --token 03em5o.agjpy9wbj98izbpn \
    --discovery-token-ca-cert-hash
sha256:92a1cf1e452961a550f05eff32f25500fc9bcfd5d5f771689e8ee2
1b6a0da243

#Copy kubeadm join token from master output of kubeadm init
command and execute in #Worker Nodes to join to cluster.

#kubectl commands has to be executed in master machine.

#Check Nodes
#=====
kubectl get nodes

```

=====End of Worker Node Script=====

### Deploy Sample Application

=====

```

kubectl run nginx-demo --image=nginx --port=80

kubectl expose pod nginx-demo --port=80 --target-port=80 --
type=NodePort

#Get Node Port details
#=====
kubectl get services

```

```
#=====
install-kubect1.sh
curl -LO "https://dl.k8s.io/release/$(curl -L -s
https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubect1
"
curl -LO
https://dl.k8s.io/release/v1.26.0/bin/linux/amd64/kubect1
sudo install -o root -g root -m 0755 kubect1
/usr/local/bin/kubect1
chmod +x kubect1
mkdir -p ~/.local/bin
mv ./kubect1 ~/.local/bin/kubect1
# and then append (or prepend) ~/.local/bin to $PATH
kubect1 version -client
```

## How to deploy/run/execute tasks/workloads in kubernetes

In Kubernetes, deploying components (such as pods, services, or other resources) can be done using either an imperative or declarative approach. Each approach offers different ways to manage resources within a Kubernetes cluster.

### Imperative Approach

The imperative approach involves issuing direct commands to the Kubernetes API server to create, modify, or delete resources. This approach is often used for interactive or one-time tasks and is beneficial for immediate changes.

#### Characteristics:

- **Direct Actions:** Commands specify the action to be taken (e.g., create, modify, delete).
- **Immediate Changes:** Changes take effect immediately.
- **Interactive:** Suitable for one-off or immediate tasks.



### Examples:

- Create a Pod:

```
apiVersion: v1
kind: Pod
metadata:
  name: <podName>
  namespace: <namespaceName>
  labels:
    key: <value>
    key: <value>
spec:
  containers:
    - name: containerName
      image: imageName
      ports:
        - containerPort: podNumber
```

- Scale a Deployment:

```
kubectl scale deployment my-deployment --replicas=3
```

- Delete a Service:

```
kubectl delete service my-service
```

### Declarative Approach

The declarative approach involves specifying the desired state of resources using configuration files (usually in YAML format) and applying them to the Kubernetes cluster. This approach is ideal for managing resources in a consistent, reproducible, and automated way.

Characteristics:

- **Desired State:** You define the desired state of resources in a configuration file.

- **Consistent and Reproducible:** Changes are applied consistently and can be versioned and shared.
- **Automation:** Files can be managed with version control systems and automation tools.

Examples:

- Create or Update a Deployment:
  - Define the desired state in a YAML file (**deployment.yaml**):

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-container
          image: nginx
```

- Apply the configuration:

```
kubectl apply -f deployment.yaml
```

- Delete Resources:
  - Define the resources you want to delete in a YAML file (**delete.yaml**):

```
apiVersion: v1
```

```
kind: Service
metadata:
  name: my-service
```

- Apply the deletion:

```
kubectl delete -f delete.yaml
```

### Choosing an Approach

- **Imperative:** Use the imperative approach for immediate, interactive tasks such as debugging or quick fixes.
- **Declarative:** Use the declarative approach for consistent and reproducible deployments, especially in production environments. This approach aligns well with DevOps practices and automation tools.

### More about Pods

Again, pods are the smallest and most basic deployable units in kubernetes. A pod represents a single instance of an application and can contain one or more containers that share the same network and storage resources. Pods are versatile and can be managed in various ways depending on the use case and application requirements. Here are the general types of pods in Kubernetes:

#### 1. Regular Pods:

- Regular pods are the most common type of pods in Kubernetes and are managed by higher-level controllers such as Deployments, StatefulSets, or DaemonSets.
- These pods typically contain one or more containers and share the same network namespace, allowing them to communicate directly with each other using **localhost**.

#### 2. Static Pods:

- Static pods are managed directly by the kubelet on each node and are defined using local YAML configuration files stored in the **/etc/kubernetes/manifests/** directory on each node.
- They are used for critical system components, such as control plane services, and are node-specific, meaning they run only on the node where their configuration resides.

### 3. Multi-Container Pods:

- Pods can contain multiple containers that work together as a single unit. These containers share the same network namespace and can communicate directly with each other using **localhost**.
- Multi-container pods are useful when you want tightly coupled containers that collaborate on a task, such as a main application container and a sidecar container for logging or monitoring.

### 4. Debug Pods:

- Debug pods are temporary pods created for troubleshooting and debugging purposes.
- These pods may contain tools for investigating issues in other running pods, such as network connectivity or application errors.

### Pod Management Using Controllers:

- Although not strictly types of pods, the way pods are managed often depends on the controller that oversees them. The main types of controllers include:
  - The **ReplicaSet**: It manages the lifecycle of ReplicaSets, which ensure a specified number of identical pods are running at all times. It automatically creates, deletes, and scales pods to maintain the desired state defined in the ReplicaSet.

```
# Define a ReplicaSet
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset # Name of the ReplicaSet
spec:
  replicas: 3 # Desired number of replicas (pods)
  # Selector specifies how to identify pods managed
  by this ReplicaSet
  selector:
    matchLabels:
      app: my-app
  # Template specifies the pod template for creating
  pods
  template:
```

```

metadata:
  # Labels are essential for matching pods with
  the selector
  labels:
    app: my-app
spec:
  # Define the container(s) for the pods
  containers:
    - name: my-container
      image: nginx # Docker image for the
  container
      ports:
        - containerPort: 80 # Expose port 80
  within the container

```

- **Deployments:** Manage replicas of identical pods, enabling rolling updates, rollbacks, and scaling.

```

# Define a Deployment
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-deployment # Name of the deployment
spec:
  replicas: 3 # Desired number of replicas
  selector:
    # Selector that matches pods based on their
  labels
    matchLabels:
      app: my-app
  template:
    # Pod template that defines how the pods should
  be created
    metadata:

```

```

    labels:
      app: my-app    # Pod labels matching the
selector
    spec:
      containers:
        - name: my-container
          image: nginx # Docker image used for the
pod
          ports:
            - containerPort: 80    # Expose port 80
within the container

```

- **StatefulSets:** Manage stateful pods, providing ordered deployment and stable network identities for each pod.

```

# Define a StatefulSet
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: my-statefulset # Name of the StatefulSet
spec:
  serviceName: "my-service" # Name of the headless
service used for network identity
  replicas: 3 # Number of pods to create
  selector:
    matchLabels:
      app: my-app # Selector that matches pods based
on their labels
  template:
    metadata:
      labels:
        app: my-app # Pod labels matching the
selector
    spec:

```

```

    containers:
      - name: my-container
        image: nginx # Docker image used for the pods

        ports:
          - containerPort: 80 # Expose port 80 within the container

```

- **DaemonSets:** Ensure that a copy of a specific pod runs on every node (or selected nodes) in the cluster.

```

# Define a DaemonSet
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: my-daemonset # Name of the DaemonSet
spec:
  selector:
    matchLabels:
      app: my-app # Selector that matches pods based on their labels
  template:
    metadata:
      labels:
        app: my-app # Pod labels matching the selector
    spec:
      containers:
        - name: my-container
          image: nginx # Docker image used for the pods

          ports:
            - containerPort: 80 # Expose port 80 within the container

```

- **Jobs and CronJobs:** Manage pods that run specific tasks to completion (Jobs) or at scheduled intervals (CronJobs).

```
# Define a CronJob
apiVersion: batch/v1
kind: CronJob
metadata:
  name: my-cronjob  # Name of the CronJob
spec:
  # Schedule for the job using cron syntax
  schedule: "*/5 * * * *"  # Run the job every 5
minutes
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: my-container
              image: busybox  # Docker image for the
pod
              # Command to run within the container
              command: ["echo", "Hi, Kubernetes!"]
              # Restart policy for the CronJob's pod(s);
              "OnFailure" means the job restarts if it fails
          restartPolicy: OnFailure
```

#### **Pod Lifecycle and Communication:**

- Pods have a lifecycle that includes phases such as:
  - Pending
  - Running
  - Terminating.



- Pods can communicate with each other directly using IP addresses and ports or through services that provide stable network endpoints for a group of pods.

kubectl get pod:

This command by default will search for pods in the default namespace which is created by default.

The above command can be run while explicitly specify the intended namespace with “-n” option in the command as thus:

```
kubectl get pod -n dev
```

Where “dev” is the intended namespace.

To completely change the default namespace context, run the command below:

```
kubectl config set-context --current --namespace=dev
```

The above command will change the default namespace context to “dev” namespace

### **How do pods communicate with each other?**

In Kubernetes, pods can communicate with each other even when they are running on different nodes. This is achieved through the cluster's networking model, which allows every pod to have a unique IP address, and through services, which provide a stable network endpoint for accessing a group of pods. Below are explanations and code examples demonstrating how communication between different pods works in a Kubernetes cluster.

#### **1. Direct Communication Using Pod IPs:**

- **Get Pod IP Addresses:**
  - Use **kubectl get pods -o wide** to list the IP addresses of pods.

```
kubectl get pods -o wide
```

- **Direct Communication:**

- Once you have the IP addresses of the pods, you can communicate directly between them using the IP and a specific port.
- For example, if a pod with IP address **10.0.0.1** has a service running on port **8080**, another pod can communicate with it by sending a request to

```
http://10.0.0.1:8080.
```

## 2. Communication via Services:

- **Create a Service:**

- Services provide an abstraction layer over a set of pods, allowing other pods to communicate with them using a stable network endpoint.
- Define a service in a YAML file (**service.yaml**):

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: my-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
```

- Apply the service:

```
kubectl apply -f service.yaml
```

- **Communicate with the Service:**

- Once the service is created, other pods can communicate with it using its DNS name (e.g., **my-service**) or its Cluster IP.
- Example command (from within another pod):

```
curl http://my-service
```

### 3. Ingress for External Communication:

- **Ingress:**

- Ingress provides a way to expose services externally with features such as load balancing, SSL termination, and path-based routing.
- Define an Ingress in a YAML file (**ingress.yaml**):

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: my-ingress
spec:
  rules:
  - http:
      paths:
      - path: /my-app
        pathType: Prefix
        backend:
          service:
            name: my-service
            port:
              number: 80
```

- Apply the Ingress:

```
kubectl apply -f ingress.yaml
```

- **Accessing a Pod via Ingress:**

- Once the Ingress is set up, you can access the pods within **my-service** using the Ingress controller's external IP and the specified path.

**Considerations:**

- **Network Plugins:**

- Kubernetes uses network plugins (e.g., Calico, Flannel, Weave) to provide networking capabilities. These plugins ensure connectivity and can impact how pods communicate with each other.

- **Network Policies:**

- You can use Network Policies to control and restrict communication between pods based on specific criteria such as labels, namespaces, or IP ranges.

**Some commonly used kubectl commands**

Kubernetes provides a powerful command-line tool called **kubectl** that allows you to manage different components within a Kubernetes cluster. Below is a table of common **kubectl** commands for managing various Kubernetes resources, including pods, deployments, services, and more:

Resource	Action	Command
Pods	List all pods	kubectl get pods
	Get details of a specific pod	kubectl describe pod <pod-name>
	Create a pod from a YAML file	kubectl apply -f <pod-file.yaml>
	Delete a pod	kubectl delete pod <pod-name>
	Stream logs from a pod	kubectl logs <pod-name>
	Exec into a pod	kubectl exec -it <pod-name> -- <command>

	Port forward to a pod	kubectl port-forward <pod-name> <local-port>:<pod-port>
Deployments	List all deployments	kubectl get deployments
	Get details of a specific deployment	kubectl describe deployment <deployment-name>
	Create a deployment from a YAML file	kubectl apply -f <deployment-file.yaml>
	Scale a deployment	kubectl scale deployment <deployment-name> --replicas=<count>
	Roll back a deployment	kubectl rollout undo deployment/<deployment-name>
	Monitor deployment status	kubectl rollout status deployment/<deployment-name>
Services	List all services	kubectl get services
	Get details of a specific service	kubectl describe service <service-name>
	Create a service from a YAML file	kubectl apply -f <service-file.yaml>
	Delete a service	kubectl delete service <service-name>
Ingress	List all ingress resources	kubectl get ingress
	Create an ingress from a YAML file	kubectl apply -f <ingress-file.yaml>
	Delete an ingress	kubectl delete ingress <ingress-name>
ConfigMaps	List all config maps	kubectl get configmaps
	Create a config map from a file	kubectl create configmap <configmap-name> --from-file=<filename>
	Create a config map from YAML	kubectl apply -f <configmap-file.yaml>
	Get details of a specific config map	kubectl describe configmap <configmap-name>
	Delete a config map	kubectl delete configmap <configmap-name>
Secrets	List all secrets	kubectl get secrets

	Create a secret from a file	kubectl create secret generic <secret-name> -from-file=<filename>
	Create a secret from YAML	kubectl apply -f <secret-file.yaml>
	Get details of a specific secret	kubectl describe secret <secret-name>
	Delete a secret	kubectl delete secret <secret-name>
Namespaces	List all namespaces	kubectl get namespaces
	Create a namespace	kubectl create namespace <namespace-name>
	Delete a namespace	kubectl delete namespace <namespace-name>
Node Management	List all nodes	kubectl get nodes
	Get details of a specific node	kubectl describe node <node-name>
	Cordon a node	kubectl cordon <node-name>
	Uncordon a node	kubectl uncordon <node-name>
	Drain a node	kubectl drain <node-name> --ignore-daemonsets --delete-local-data
DaemonSets	List all daemonsets	kubectl get daemonsets
	Create a daemonset from YAML	kubectl apply -f <daemonset-file.yaml>
	Delete a daemonset	kubectl delete daemonset <daemonset-name>
StatefulSets	List all statefulsets	kubectl get statefulsets
	Create a statefulset from YAML	kubectl apply -f <statefulset-file.yaml>
	Delete a statefulset	kubectl delete statefulset <statefulset-name>
Jobs	List all jobs	kubectl get jobs
	Create a job from YAML	kubectl apply -f <job-file.yaml>
	Delete a job	kubectl delete job <job-name>
CronJobs	List all cronjobs	kubectl get cronjobs
	Create a cronjob from YAML	kubectl apply -f <cronjob-file.yaml>
	Delete a cronjob	kubectl delete cronjob <cronjob-name>

## Differences between ReplicaSets and Replication Controller

ReplicaSet	Replication Controller
Supports equality-based selectors and also set based selectors.	Only support equality-based selectors
Supper set of replication controller	Subset of ReplicaSet
Equality based example: Key == value(Equality condition)	Set based example: Key in [ value1, value2, value3] Selector: matchLabels: <i># Equality Based</i> key: value matchExpressions: <i># Set Based</i> - Key:app Operator: in Values: - javaWebApp - myApp -fe

### ReplicaSet Example:

```
Kind: ReplicaSet
apiVersion: apps/v1
metadata:
  name: <RSName>
  labels:
    <key>: <value>
spec:
  selector:
    matchLabels: # Equality Based
    <key>: <value>
    matchExpressions:
      - key: <key>
        operator: <in / not in>
        values:
          - <value1>
          - <value2>
          - <value3>
  replicas: <RSNumber>
  template:
    metadata:
      name: <podName>
      labels:
        <key>: <values>
    spec:
      containers:
        - name: <containerName>
          image: <imageName:tag>
          ports:
            - containerPort: <ContainerPodNumber>
```

### Example: Deploy a nodeApp in a ReplicaSet

```
kind: ReplicaSet
apiVersion: apps/v1
metadata:
  name: nodeApp
spec:
  replicas: 2
  selector:
    matchLabels: # Equality Based
      app: node
  template:
    metadata:
      name: nodeWebApp
      labels:
        app: node
    spec:
      imagePullSecrets:
        - name: dockerhublogin
      containers:
        - name: nodeApp
          image: fewaitconsulting/nodejs-fe-app
          ports:
            - containerPort: 9981

# service to discover the ReplicaSet
---
apiVersion: v1
kind: Service
metadata:
  namespace: dev
  name: nodeAppsvc
spec:
  selector:
    app: node
    type: nodePort
  ports:
    targetPort: 8080
    nodePort: 31500 # Range of node ports 30000 - 32676
    port: 80
```

**NB:** When pulling the docker image from a private docker registry, Kubernetes will need to be authenticated so a secret need to be created and its reference include the above file.

```
kubectrl create secret docker-registry dockerhublogin \
--docker-server=docker.io \
--docker-username=DOCKER_USER \
--docker-password=DOCKER_PASSWORD \
--docker-email=DOCKER_EMAIL
```





**NB:** To access this application on the browser, we need to use the nodePort address together with the server IP address.

## DaemonSet:

It is almost the same as ReplicaSet apart of the fact that they don't have number of replicas. DaemonSet automatically schedule a copy of the same pod for all nodes or groups of nodes.

### Example: Deploy a log management (logmgmt) in a DaemonSet.

```
kind: DaemonSet
apiVersion: apps/v1
metadata:
  name: logmgmt
spec:
  selector:
    matchLabels: # Equality Based
    app: node
  template:
    metadata:
      name: hello
      labels:
        app: hello
    spec:
      imagePullSecrets:
        - name: dockerhublogin
      containers:
        - name: hello
          image: fewaitconsulting/hello
          ports:
            - containerPort: 9981

# service to discover the ReplicaSet
---
apiVersion: v1
kind: Service
metadata:
  namespace: dev
  name: nodeAppsvc
spec:
  selector:
    app: hello
    type: nodePort
  ports:
    targetPort: 8080
    nodePort: 31500 # Range of node ports 30000 - 32676
    port: 80
```

NB: The DaemonSet will schedule the application in the other nodes except the master node which is by default tainted i.e. The node doesn't accept application scheduling in it.

A node is tainted if it needs:

- Recommissioning
- Upgrades
- Updating
- Patching

How to taint a node:

```
kubectl taint nodes node1 key1=value1:NoSchedule
```

How to untaint a node:

```
kubectl taint nodes node1 key1=value1:NoSchedule-
```

NB: An application can be scheduled on all the nodes whether tainted or not by adding the tolerations to the daemonSet manifest file. Example: Let us modify the above daemonSet file to do that:

```
kind: DaemonSet
apiVersion: apps/v1
metadata:
  name: logmgmt
spec:
  selector:
    matchLabels: # Equality Based
      app: node
  template:
    metadata:
      name: hello
      labels:
        app: hello
    spec:
      imagePullSecrets:
        - name: dockerhublogin
      containers:
        - name: hello
          image: fewaitconsulting/hello
          ports:
            - containerPort: 9981
      tolerations:
        operator: Exists
        effect: "NoSchedule"
```

## Deployment controller

Deployment will exhibit the following characteristics:

- Deploy a ReplicatSet (RS)
- Update pods (PodTemplateSpec)
- Rollback to older deployment versions
- Scale deployment up or down
- Pause and resume deployment.
- Use the status of the deployment to determine state of replicas
- Clean up older RS that aren't needed anymore

Deployment has two strategies:

### 1. Recreate (Blue-Green) Deployment:

- In this strategy, a new version of the application is deployed by replacing the existing version entirely.
- The old version (blue) is shut down, and the new version (green) is started.
- While simple, it may cause downtime during the switch from the old to the new version.

### 2. Rolling Update Deployment:

- Rolling updates gradually replace instances of the old version with instances of the new version, one at a time.
- Pods are replaced gradually, ensuring that the application remains available throughout the update process.
- Rollback is straightforward if any issues are detected during the update.

**Example: Deployment of a hello app in a deployment using Recreate strategy.**

```
kind: Deployment
apiVersion: apps/v1
metadata:
  name: hello
spec:
  replicas: 2
  selector:
    matchLabels: # Equality Based
```

```

    app: hello
  strategy:
    type: recreate
  template:
    metadata:
      name: hellopod
      labels:
        app: hello
    spec:
      imagePullSecrets:
        - name: dockerhublogin
      containers:
        - name: helloworld
          image: fewaitconsulting/hello-app
          ports:
            - containerPort: 80

# service to discover the Deployment
---
apiVersion: v1
kind: Service
metadata:
  namespace: dev
  name: hellosvc
spec:
  selector:
    app: hello
    type: nodePort
  ports:
    targetPort: 80
    nodePort: 31000 # Range of node ports 30000 - 32676
    port: 80

```

## Example 2: Deployment of javawebapp application using RollingUpdate strategy

```

kind: Deployment
apiVersion: apps/v1
metadata:
  name: javawebapp
spec:
  replicas: 2
  selector:
    matchLabels: # Equality Based
      app: javawebapp
  strategy:
    type: RollingUpdate
    maxUnavailable: 1
    maxSurge: 1
    minReadySeconds: 30
  template:
    metadata:
      name: javawebapppod
      labels:

```

```

        app: javawebapp
    spec:
        imagePullSecrets:
            - name: dockerhublogin
        containers:
            - name: helloworld
              image: fewaitconsulting/java-web-app
              ports:
                  - containerPort: 8080


# service to discover the Deployment
---
apiVersion: v1
kind: Service
metadata:
    namespace: dev
    name: javawebappsvc
spec:
    selector:
        app: javawebapp
        type: nodePort
    ports:
        targetPort: 80
        nodePort: 31000 # Range of node ports 30000 - 32676
        port: 80

```

## Deployment techniques:

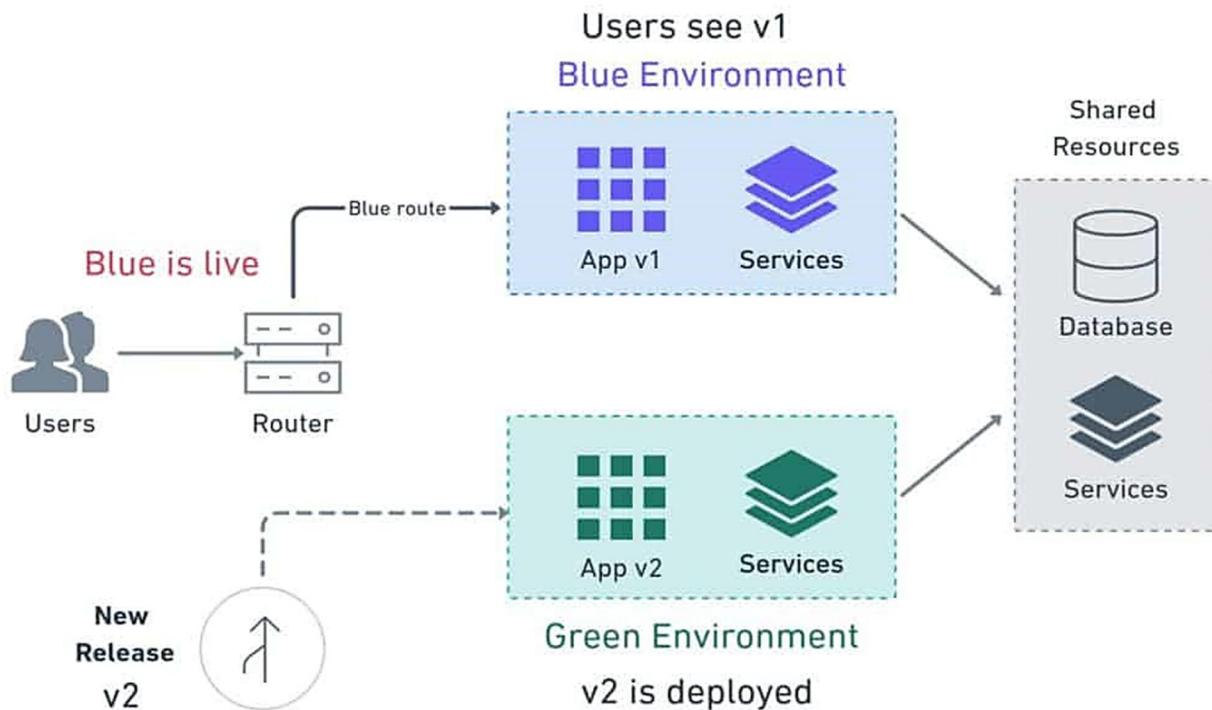
### 1. Canary Deployments:

- Canary deployments involve gradually rolling out a new version of an application to a subset of users or traffic while monitoring its performance and stability.
- Traffic routing mechanisms, such as Kubernetes Ingress controllers or service meshes, are used to direct a portion of the traffic to the new version.
- If the new version performs well, more traffic is gradually shifted to it. Otherwise, the deployment can be rolled back.



## 2. Blue-Green Deployments:

- Blue-green deployments involve maintaining two identical environments (blue and green) running simultaneously, with only one environment serving live traffic at a time.
- New versions of the application are deployed to the inactive environment, and traffic is switched from the active environment to the updated one once the deployment is successful.
- This approach enables zero-downtime deployments and allows for easy rollback in case of issues.



### Resource Request and Limit

- Resource requests define the minimum amount of CPU and memory that a container needs to run.
- Kubernetes scheduler uses resource requests to decide where to place pods in the cluster based on available resources.
- Limit is the maximum number of resources that Kubernetes will allow the container to use.

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
  - name: my-container
    image: nginx
    resources:
      requests: # Resource requests specify the minimum required
                 resources for the container
                 cpu: "100m" # 100 millicpu (0.1 CPU)
                 memory: "128Mi" # 128 Mebibytes
```



```
limits: # Resource limits specify the maximum allowed resources
for the container
cpu: "200m" # 200 millicpu (0.2 CPU)
memory: "256Mi" # 256 Mebibytes
```

## Pod Auto-Scaling

Pod autoscaling, also known as **Horizontal Pod Autoscaler** (HPA), is a Kubernetes feature that dynamically adjusts the number of pod replicas based on observed resource utilization metrics, such as CPU and memory usage.

The HorizontalPodAutoscaler controller continuously monitors these metrics and scales the number of pod replicas up or down to meet specified target resource utilization levels.

Pod autoscaling helps ensure optimal performance, availability, and cost efficiency by automatically adapting the infrastructure to changing workloads.

## Differences between Pod Auto-scaling and AWS Auto-scaling

Aspect	Pod Autoscaling (Kubernetes)	AWS Autoscaling
<b>Scope</b>	Operates at the level of individual pods	Operates at the level of EC2 instances or other AWS resources
<b>Managed Resources</b>	Manages scaling of pod replicas within Kubernetes cluster	Manages scaling of EC2 instances or other AWS resources within AWS environment
<b>Resource Metrics</b>	Scales pods based on resource utilization metrics such as CPU and memory	Can scale resources based on various metrics including CPU utilization, network traffic, and custom CloudWatch metrics
<b>Integration</b>	Integrated with Kubernetes' native scaling mechanisms (e.g., Horizontal Pod Autoscaler)	Integrated with various AWS services including EC2, ECS, and ELB
<b>Environment</b>	Typically used in containerized environments managed by Kubernetes	Used within AWS cloud environment to manage scalability of AWS resources

## How to check resource usage in Kubernetes

### 1. **kubectl top:**

- Use the **kubectl top** command to view resource usage of pods, nodes, or namespaces.
- For example, to view CPU and memory usage of pods:

```
kubectl top pods
```

- To view CPU and memory usage of nodes:

```
kubectl top nodes
```

**NB:** The two above commands require the Metrics Server to be installed in your cluster.

### 2. **Dashboard:**

- Kubernetes Dashboard provides a web-based UI for monitoring cluster resources.
- You can view resource usage metrics, including CPU and memory, for pods, nodes, and namespaces.

### 3. **Metrics Server:**

- The Metrics Server collects resource usage metrics from pods and nodes in the cluster.
- You can access metrics exposed by the Metrics Server through the Kubernetes API.

### 4. **Prometheus:**

- Prometheus is a popular monitoring and alerting tool used with Kubernetes.
- You can configure Prometheus to scrape metrics from Kubernetes components and applications, including resource usage metrics.

### 5. **Third-party Monitoring Solutions:**

- There are many third-party monitoring solutions available for Kubernetes, such as Datadog, Grafana, and Prometheus Operator.
- These tools offer advanced monitoring capabilities and customizable dashboards for visualizing resource usage metrics.

### Installation of Metric Server

1. **Download Metric Server YAML files:** You can clone the Metric Server repository from GitHub:

```
wget https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
```

2. **Deploy the Metric Server:** Apply the YAML files to deploy the Metric Server components:

```
kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
```

3. **Verify the deployment:** Check the status of Metric Server components to ensure they are running:

```
kubectl get pods -n kube-system
```

You should see the **metrics-server** deployment running in the **kube-system** namespace.

4. **Verify Metric Server metrics:** Check if Metric Server metrics are available:

```
kubectl top nodes
```

If Metric Server is correctly installed and configured, you should see resource usage metrics for nodes in the cluster.

## Deployment with HPA

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hpadeployment
spec:
  selector:
    matchLabels:
      name: hpapod
  template:
    metadata:
      labels:
        name: hpapod
    spec:
      containers:
        - name: hpacontainer
          image: registry.k8s.io/hpa-example
          ports:
            - containerPort: 80
          resources:
            limits:
              cpu: "100m"
              memory: "64Mi"
            requests:
              cpu: 200m
              memory: "256Mi"
```

```
---
apiVersion: v1
kind: Service
metadata:
  name: hpacusterservice
  labels:
    name: hpaservice
spec:
  ports:
    - port: 80
  selector:
    name: hpapod
  type: NodePort
```

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: hpadeploymentautoscaler
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: hpadeployment
  minReplicas: 2
  maxReplicas: 10
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 50
    - type: Resource
      resource:
        name: memory
```

```
target:
  type: Utilization
  averageUtilization: 50
```

### Example: Deployment of Spring app and MongoDB as pod without volumes

#### Requirements:

- Stateless application deployment.
- Deployment is the best choice Kubernetes object
- MongoDB:
  - ✓ Use StatefulSet as the choice Kubernetes object since it is stateful application.

```
# mongodb-replicaset.yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: mongodb
  namespace: mongodb
spec:
  replicas: 3
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: mongodb
    spec:
      containers:
        - name: mongodbcontainer
          image: mongo:4.2
          ports:
            - containerPort: 27017
          volumeMounts:
            - name: mongodb-storage
              mountPath: /data/db
          env:
            name: MONGO_INITDB_ROOT_USERNAME: admin
            value: MONGO_INITDB_ROOT_PASSWORD: admin123
            name: MONGO_INITDB_DATABASE: mydatabase
            value: devdb@123
---
# mongodb-service.yaml
apiVersion: v1
kind: Service
metadata:
  name: mongosvc
spec:
  type: ClusterIP
  selector:
```

```

    app: mongo
    ports:
      port: 27017
      targetPort: 27017
---
# springapp-service.yaml
apiVersion: v1
kind: Service
metadata:
  name: springappsvc
spec:
  type: NodePort
  selector:
    app: springapp
  ports:
    port: 80
    targetPort: 8080

```

## Kubernetes Volumes:

### HostPath Volume:

```

apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: mongodb
spec:
  replicas: 3
  selector:
    matchLabels:
      app: mongo
  template:
    metadata:
      name: myapp
      labels:
        app: mongo
    spec:
      containers:
        volumes:
          - name: mongodbbhostvol
            hostpath:
              path: /mongodata
          - name: vol2
            hostPath:
              path: /tmp/data
      containers:
        - name: mongodbbcontainer
          image: mongo
          ports:
            - containerPort: 27017
          env:
            - name: MONGO_INITDB_ROOT_USERNAME: admin
              value: MONGO_INITDB_ROOT_PASSWORD: admin123

```

```
- name: MONGO_INITDB_DATABASE: mydatabase
  value: devdb@123
volumeMounts:
- name: mongodbdhostvol
  mountPath: /data/db
- name: vol2
  mountPath: /data/db
```

## Configuration of NFS Server

Step 1: Create one server for NFS

NFS = Network File System. It is a distributed file system shared file system.

Update our system's repository index that of the internet through the following apt command as sudo:

```
sudo apt-get update
```

The above command let us install the latest available version of a software through the ubuntu repositories.

Now, run the following command in order to install the NFS kernel server on your system:

```
sudo apt install nfs-kernel-server -y
```

Step 2: Create the export directory

```
sudo mkdir -p /mnt/share
```

**NB:** As we want all clients to access the directory, we will remove restrictive permissions.

```
sudo chown nobody:nogroup /mnt/share
sudo chmod 777 /mnt/share
```

Step 3: Assign server access to client(s) through NFS export file

```
sudo vi /etc/exports
```

```
# /mnt/share <YourClientIP or clients>(rw,sync,no_subtree_check) OR
# /mnt/share *(rw,sync,no_subtree_check) # To permit all IP addresses
```

```
/mnt/share 192.168.1.0/24(rw,sync,no_subtree_check)
```

#### Step 4: Export the shared directory

```
sudo exportfs -a  
sudo systemctl restart nfs-kernel-server
```

#### Step 5: Open firewall for the client (s) port 2049

### Configuring the NFS Client Machine

Step 1: Install NFS common on all the worker nodes that need the NFS server.

Install the NFS client package

```
sudo apt update  
sudo apt install nfs-common -y
```

### Deployment of a MongoDB using the NFS

```
# mongo-nfs.yaml  
apiVersion: apps/v1  
kind: ReplicaSet  
metadata:  
  name: mongodb  
spec:  
  replicas: 3  
  selector:  
    matchLabels:  
      app: mongo  
  template:  
    metadata:  
      name: myapp  
      labels:  
        app: mongo  
    spec:  
      containers:  
        volumes:  
          - name: mongodbvols  
            nfs:  
              server: #Your server IP  
              path: /mnt/share  
      containers:  
        - name: mongodbcontainer  
          image: mongo  
          ports:  
            - containerPort: 27017  
          env:  
            - name: MONGO_INITDB_ROOT_USERNAME: admin  
              value: MONGO_INITDB_ROOT_PASSWORD: admin123
```



```
- name: MONGO_INITDB_DATABASE: mydatabase
  value: devdb@123
volumeMounts:
- name: mongodbv1
  mountPath: /data/db
- name: vol2
  mountPath: /data/db
```

## Setting up Kubernetes cluster on AWS using KOPS

- It is a software use to create production ready k8s cluster in a cloud provider like AWS
- It supports multiple cloud providers
- It competes with managed k8s services like EKS, AKS and GKE
- It is cheaper than the others
- It creates production ready k8s
- It creates resources like: LoadBalancer, AutoScaling Groups, Launch configuration, worker node, master node (Control plane)
- It is IaC

```
#!/bin/bash
```

1) Create Ubuntu EC2 instance in AWS

2) a) Create kops user

```
sudo adduser kops
sudo echo "kops ALL=(ALL) NOPASSWD:ALL" | sudo tee /etc/sudoers.d/kops
sudo su - kops
```

Install AWSCLI using the apt package manager

```
sudo apt install awscli -y
```

OR

2b) Install AWSCLI using the script below

```
sudo apt update -y
sudo apt install wget unzip -y
```

```
sudo curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o
"awscliv2.zip"
sudo apt install unzip python -y
sudo unzip awscliv2.zip
sudo ./awscli-bundle/install -i /usr/local/bin/aws -b /usr/local/bin/aws
```

3) Install kops software on an ubuntu instance by running the commands below:

```
sudo apt install wget -y
sudo wget https://github.com/kubernetes/kops/releases/download/v1.22.0/kops-
linux-amd64
chmod +x kops-linux-amd64
sudo mv kops-linux-amd64 /usr/local/bin/kops
```

4) Install kubectl kubernetes client if it is not already installed

```
curl -LO curl -LO https://dl.k8s.io/release/v1.30.0/bin/linux/arm64/kubectl

chmod +x ./kubectl
sudo mv ./kubectl /usr/local/bin/kubectl
```

5) Create an IAM role from AWS console or CLI with the below policies

- ✓ AmazonEC2FullAccess
- ✓ AmazonS3FullAccess
- ✓ IAMFullAccess
- ✓ AmazonVPCFullAccess

Then attach IAM role to ubuntu server from the server console. Select created KOPS server→Actions→Instance Setting→Attach/Replace IAM Role→Select the role which you created→Save

6) Create an S3 bucket

Execute the commands below in your KOPS control server. Use unique S3 bucket name, if you get a bucket name exist error.

```
aws s3 mb s3://fewaitconsulting
aws s3 ls # To verify if s3 bucket is created
```

7) Expose environmental variables of the created bucket

```
#Add env variables in bashrc

sudo vi .bashrc

#!/bin/bash
```

```
# Give unique name to the created s3 bucket
export NAME=fewaitconsulting.k8s.local
export KOPS_STATE_STORE=s3://fewaitconsultingkops

source ~/.bashrc
```

#### 8) Create ssh keys before creating cluster

```
ssh-keygen
```

#### 9) Create kubernetes cluster definitions on s3 bucket

```
kops create cluster --zones us-east-1a --networking weave --master-size
t2.medium --master-count 1 --node-size t2.medium --node-count=2 ${NAME}
#Copy your sshkey into your cluster to be able to access your kubernetes
nodes
kops create secret --name ${NAME} sshpublickey admin -i ~/.ssh/id_rsa.pub
```

#### 10) Initialize your kops kubernetes cluster by running the command below

```
kops update cluster {NAME} --yes
```

#### 11) Validate your cluster (KOPS will take some time to create cluster. Execute commands below after 3 or 4 mins)

```
kops validate cluster
```

OR

Export the kubeconfig file to manage your kubernetes cluster from a remote server. For this demo, our remote server shall be our kops server.

```
kops export kubecfg $NAME --admin
```

#### 12) To list nodes and pods to ensure that you can make calls to the kubernetes apiServer and run workloads, use the command below:

```
kubectl get nodes

#Number and stopping of nodes can be altered in the autoscaling group
service back in AWS
```

## Persistent Volumes

- It is a piece of storage (hostPath, nfs, ebs, azurefile, azuredisk) in k8s cluster
- It exists independently from pod life cycle from which it is consuming

## PersistentVolumeClaim

- It requests for storage (volume).
- Using PVC, we can request (specify) how much storage we need and with what access mode we need the volume.
- Persistent volumes are provisioned in two ways: Statistically or Dynamically.

### 1) Static volumes (manual provisioning)

- A k8s administrator can create a PV manually so that PV can be available for pods which requires.
- Create a PVC so that PVC will attach PV. We can use PVC with pods to get an access to PV.

### 2) Dynamic volumes (Dynamic provisioning)

- It's possible to have k8s provision (create) volumes (PV) as required, provided we have configured storage.

So, when we create PVC, if PV is not available, storage class will create PV dynamically.

PVC: If pod requires access to storage (PV), it will get an access using PVC. PVC will be attached to PV.

- ✓ **Persistentvolume:** The low-level representation of a storage volume.
- ✓ **PeristentVolumeClaim:** The binding between a pod and persistentvolume
- ✓ **Pod:** A running container that will consume a persistentvolume.
- ✓ **StorageClass:** It allows for dynamic provisioning of persistentvolumes.

PersistentVolume will have access modes:

- ReadWriteOnce → The volume can be mounted as read-write by a single mode
- ReadOnlyMany → The volume can be mounted as read-only by many modes
- ReadWriteMany → The volume can be mounted as read-write by many nodes

## Claim Policies

A persistent volume claim can have several different claims policies associated with it including:

- Retain→When the claim (PVC) is deleted, the volume (PV) will exist
- Recycle→ When the claim is deleted, then the volume remains but, in a state, where the data can be manually recovered.
- Delete→ The persistent volume is deleted when the claim is deleted.

The claim policy (associated with the PV and not the PVC) is responsible for what happens to the data when the claim is deleted.

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-hostpath
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteMany
  nfs:
    server: <nfs server ip>
    path: "/mnt/share"

---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pvc-nfs-pv1
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 500Mi
```

## Example: Deployment of MongoDB application with the use of PVC

```
# pv-pvc-mongo-hostpath.yaml

apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-hostpath
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/kube"

---

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pv-hostpath
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 300Mi
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 500Mi
  hostPath:
    path: "/kube"

---

apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: mongodbrs
spec:
  replicas: 3
  selector:
    matchLabels:
      app: mongodb
  template:
    metadata:
      name: mongodbpod
      labels:
```

```

    app: mongodb
  spec:
    containers:
      volumes:
        - name: mongodbpvc
          persistentVolumeClaim:
            claimName: pvc-hostpath
    containers:
      - name: mongodbcontainer
        image: mongodb
        ports:
          - containerPort: 27017
        env:
          - name: MONGO_INITDB_ROOT_USERNAME: admin
            value: MONGO_INITDB_ROOT_PASSWORD: admin123
          - name: MONGO_INITDB_DATABASE: mydatabase
            value: devdb@123
        volumeMounts:
          - name: mongodbv1
            mountPath: /data/db
          - name: vol2
            mountPath: /data/db

```

```

pv-pvc-mongo-nfs.yaml
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-nfs-pv1
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteMany
  nfs:
    server: <nfs server ip>
    path: "/mnt/share"
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pvc-nfs-pv1
spec:
  storageClassName: manual
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 500Mi
---
apiVersion: apps/v1

```

```

kind: ReplicaSet
metadata:
  name: mongodbrs
spec:
  selector:
    matchLabels:
      app: mongodb
  template:
    metadata:
      name: mongodbpod
      labels:
        app: mongodb
    spec:
      volumes:
        - name: mongodb-pvc
          persistentVolumeClaim:
            claimName: pvc-nfs-pv1
      containers:
        - name: mongodbcontainer
          image: mongo
          ports:
            - containerPort: 27017
          env:
            - name: MONGO_INITDB_ROOT_USERNAME
              value: devdb
            - name: MONGO_INITDB_ROOT_PASSWORD
              value: devdb@123
          volumeMounts:
            - name: mongodb-pvc
              mountPath: /data/db

```

**Complete Manifest Where in single yml we defined Deployment & Service for SpringApp & PVC (with default StorageClass),ReplicaSet & Service For Mongo.**

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: springappdeployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: springapp
  template:
    metadata:
      name: springapppod
      labels:
        app: springapp
    spec:
      containers:
        - name: springappcontainer
          image: fewaitconsulting/spring-boot-mongo

```



```

    ports:
      - containerPort: 8080
    env:
      - name: MONGO_DB_USERNAME
        value: devdb
      - name: MONGO_DB_PASSWORD
        value: devdb@123
      - name: MONGO_DB_HOSTNAME
        value: mongo
  ---

```

```

apiVersion: v1
kind: Service
metadata:
  name: springapp
spec:
  selector:
    app: springapp
  ports:
    - port: 80
      targetPort: 8080
  type: LoadBalancer
  ---

```

```

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mongodbpvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
  ---

```

```

apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: mongodbrs
spec:
  selector:
    matchLabels:
      app: mongodb
  template:
    metadata:
      name: mongodbpod
      labels:
        app: mongodb
    spec:
      volumes:
        - name: pvc
          persistentVolumeClaim:
            claimName: mongodbpvc
      containers:
        - name: mongodbcontainer
          image: mongo
          ports:
            - containerPort: 27017
          env:

```

```

- name: MONGO_INITDB_ROOT_USERNAME
  value: devdb
- name: MONGO_INITDB_ROOT_PASSWORD
  value: devdb@123
volumeMounts:
- name: pvc
  mountPath: /data/db

```

```

---
apiVersion: v1
kind: Service
metadata:
  name: mongo
spec:
  type: ClusterIP
  selector:
    app: mongoddb
  ports:
  - port: 27017
    targetPort: 27017

```

## Kubernetes ConfigMaps and Secrets

- **ConfigMaps** are used to pass additional configurations needed to orchestrate containers including environmental variables. Examples is hostname or password.
- **ConfigMaps** pass values in plain text key: value pairs
- **Kubernetes Secrets** let you store and manage sensitive information, such as passwords, OAuth tokens, and ssh keys.
- Storing confidential information in a Secret is safer and more flexible than putting it directly in a Pod definition or in a container image.

```

kind: ConfigMap
apiVersion: v1
metadata:
  name: mongo-configmap
data:
  # Configuration values can be set as key-value properties
  db-password: devdb@123
  db-hostname: devdb
---
apiVersion: v1
kind: Secret
metadata:
  name: mongo-db-password
#type: Opaque means that from kubernetes's point of view the contents of
this Secret is unstructured.
#It can contain arbitrary key-value pairs.
type: Opaque

```

```
data:
  # Output of "echo -n 'devdb@123' | base64"
  db-password: ZGV2ZGJAMTIz
```

### Example: Using of ConfigMaps and Secrets to deploy a springapp which uses a mongodb database

```
# Complete Manifest Where in single yml we defined Deployment & Service
for SpringApp & PVC(with default StorageClass),ReplicaSet & Service For
Mongo.
apiVersion: apps/v1
kind: Deployment
metadata:
  name: springappdeployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: springapp
  template:
    metadata:
      name: springapppod
      labels:
        app: springapp
    spec:
      containers:
        - name: springappcontainer
          image: mylandmarktech/spring-boot-mongo
          ports:
            - containerPort: 8080
          env:
            - name: MONGO_DB_USERNAME
              valueFrom:
                configMapKeyRef:
                  name: springappconfig
                  key: db-username
            - name: MONGO_DB_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: springappsecret
                  key: db-password
            - name: MONGO_DB_HOSTNAME
              valueFrom:
                configMapKeyRef:
                  name: springappconfig
                  key: db-hostname
---
apiVersion: v1
kind: Service
```

```

metadata:
  name: springapp
spec:
  selector:
    app: springapp
  ports:
    - port: 80
      targetPort: 8080
  type: LoadBalancer
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mongodbpvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
---
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: mongodbrs
spec:
  selector:
    matchLabels:
      app: mongodb
  template:
    metadata:
      name: mongodbpod
      labels:
        app: mongodb
    spec:
      volumes:
        - name: pvc
          persistentVolumeClaim:
            claimName: mongodbpvc
      containers:
        - name: mongodbcontainer
          image: mongo
          ports:
            - containerPort: 27017
          env:
            - name: MONGO_INITDB_ROOT_USERNAME
              value: devdb
            - name: MONGO_INITDB_ROOT_PASSWORD
              value: devdb@123
          volumeMounts:
            - name: pvc
              mountPath: /data/db
---
apiVersion: v1
kind: Service
metadata:
  name: mongo

```

```
spec:
  type: ClusterIP
  selector:
    app: mongodb
  ports:
    - port: 27017
      targetPort: 27017
```